

Advanced P2P with RTMFP: Tips and Tricks Michael Thornburgh | Senior Computer Scientist



RTMFP Refresher



- "rtmfp:"vs"rtmp:"
 - Add an "f"
- NetConnection, NetStream, NetGroup
- Client-server, 1:1 P2P, group P2P, IP multicast
- UDP vs TCP
 - But always congestion controlled!
- Full and partial reliability
 - Avoid retransmission to keep latency down
- Prioritization for real-time communication
 - Requires congestion control

RTMFP Refresher



- Communication requires a NetConnection
 - var netConnection:NetConnection = new NetConnection;
 ...
 netConnection.connect("rtmfp://server.example.com/...", ...);
- RTMFP server performs P2P introduction, NAT/firewall traversal
- Communication (including P2P) halts if NetConnection is closed or connection to the server is lost
- Server options
 - FMES4+, on premise or hosted on AWS EC2 by the hour
 - Codename Cirrus service: http://labs.adobe.com/technologies/cirrus/

RTMFP Refresher – P2P modes



- One-to-one direct connections
 - Publish:

```
    var ns:NetStream = new NetStream(netConnection, NetStream.DIRECT_CONNECTIONS);
    ns.publish("stream");
    Play:
    var ns:NetStream = new NetStream(netConnection, peerID);
    ns.play("stream");
```

- Peer ID is a unique 256-bit cryptographic pseudorandom identifier
 - Locally derived by each peer, not assigned by server
 - netConnection.nearID
 - Disappears when NetConnection closes

RTMFP Refresher – P2P modes



Groups

- Group specification string ("groupspec")
 - Start with "G:" to distinguish from peer IDs
 - GroupSpecifier class helps you make groupspecs
- Join
 - var ns:NetStream = new NetStream(netConnection, groupspec);
 - var ng:NetGroup = new NetGroup(netConnection, groupspec);
- P2P permission dialog
 - Groups can use upload bandwidth for other members' traffic
 - Must be accepted before groups will function
- NetConnection with or without server



- Multicast streaming (P2P, native IP, "fusion")
- Posting
- Directed Routing
- Object Replication



- Multicast streaming (P2P, native IP, "fusion")
 - netStream.publish("stream");
 - netStream.play("stream");
 - NetGroup.MulticastStream.PublishNotify/.UnpublishNotify events
- Posting
- Directed Routing
- Object Replication



- Multicast streaming (P2P, native IP, "fusion")
- Posting
 - netGroup.post(object)
 - NetGroup.Posting.Notify event
- Directed Routing
- Object Replication



- Multicast streaming (P2P, native IP, "fusion")
- Posting
- Directed Routing
 - netGroup.sendToNearest()/sendToNeighbor()/sendToAllNeighbors()
 - netGroup.receiveMode
 - NetGroupReceiveMode.EXACT (default)
 - NetGroupReceiveMode.NEAREST
 - NetGroup.SendTo.Notify event
- Object Replication



- Multicast streaming (P2P, native IP, "fusion")
- Posting
- Directed Routing
- Object Replication
 - netGroup.addHaveObjects()/addWantObjects()
 - netGroup.removeHaveObjects()/removeWantObjects()
 - netGroup.writeRequestedObject()
 - netGroup.replicationStrategy (lowest first, rarest first)
 - Events
 - NetGroup.Replication.Request
 - NetGroup.Replication.Fetch.Result
 - NetGroup.Replication.Fetch.SendNotify
 - NetGroup.Replication.Fetch.Failed

RTMFP Refresher



For a detailed review, watch Matthew's sessions from MAX 2008 & 2009

http://tv.adobe.com/watch/max-2008-develop/future-of-communication-with-rtmfp-by-matthew-kaufman/

http://tv.adobe.com/watch/max-2009-develop/p2p-on-the-flash-platform-with-rtmfp/

Serverless Mode

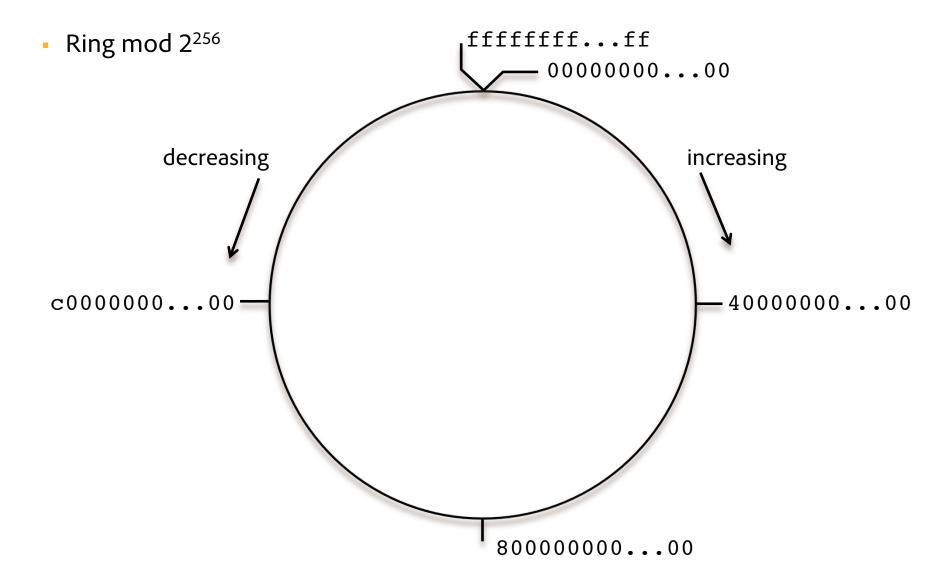


- NetConnection.connect("rtmfp:")
 - also used for receiving pure IP multicast streams
 - not talking about that today (not P2P;)
- For Groups
 - Can use 1:1 NetStreams if a low-level session is already up (via groups)
 - as of Flash Player 11.0
 - use a 2-member group to guarantee
- Must use IP multicast to discover/establish peers
 - GroupSpecifier.ipMulticastMemberUpdatesEnabled = true
 - GroupSpecifier.addIPMulticastAddress()
 - LAN only (TTL/hop limit 1)



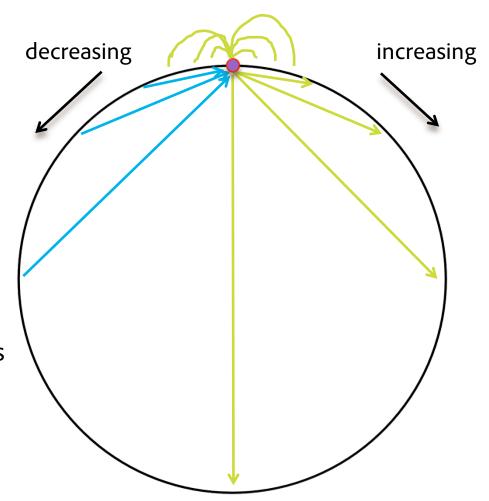
- Group address is position on ring mod 2²⁵⁶
 - SHA256 hash of peer ID (and other stuff)
 - and peer ID is itself a SHA256 hash of a cryptographic public key (and other stuff)
 - can't directly choose or influence
 - evenly distributed
 - NetGroup.estimatedMemberCount
 - NetGroup.convertPeerIDToGroupAddress()
 - represented as 64-digit hex strings (256 bits),







- 3 immediate increasing
- 3 immediate decreasing
 - maintain the ring, full transitive connectivity
- Binary fractionation ½, ¼ ...
- 6 low latency (lowest RTT)
- Binary fractionation target from other side
- 1 random every 10(ish) seconds





- Full mesh guaranteed* for groups 14 or smaller
 - 3 (inc) + 3 (dec) + 6 (fast) + 1 (rand) + 1 (self)
 - * NATs or firewalls may block some neighbor connections
 - usually not an issue for LAN communication
 - *High churn may delay convergence
 - Stale peer records may still seem worth trying for a while
 - Stale records purged after 5 minutes

Directed Routing



- GroupSpecifier.routingEnabled = true
- Send a message directly to one or more peers in a group
- Reliable, in-order delivery mode

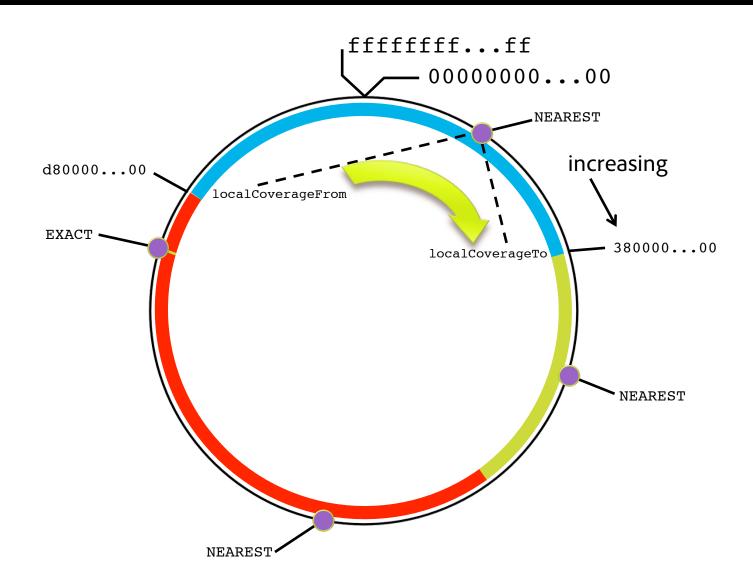
Local Coverage



- NetGroup.receiveMode
 - EXACT (default): localCoverageFrom = localCoverageTo = group address
 - NEAREST: localCoverageFrom to localCoverageTo inclusive
 - measured in increasing direction on ring
 - depends on your neighbors' addresses and receive modes
 - coverage ranges for NEAREST-mode peers in a properly connected group are contiguous and non-overlapping
 - EXACT peers may be point holes in other peers' NEAREST ranges

Local Coverage





Automatic Distributed Election



- Consistently select O(1) subset of members for special responsibilities
 - Example: Retrieve file from server and seed via Object Replication
- Local coverage can be used to implement
 - Example: peer covering "0000...00" is "elected"
 - Easy to check with string comparisons, 256-bit math not needed ©
- Gotcha when electing more than one peer:
 - Postings must have unique byte serialization
 - Consistent serialization will avoid duplicates in the network
 - Use Array, avoid Object

Automatic Distributed Election



DEMO

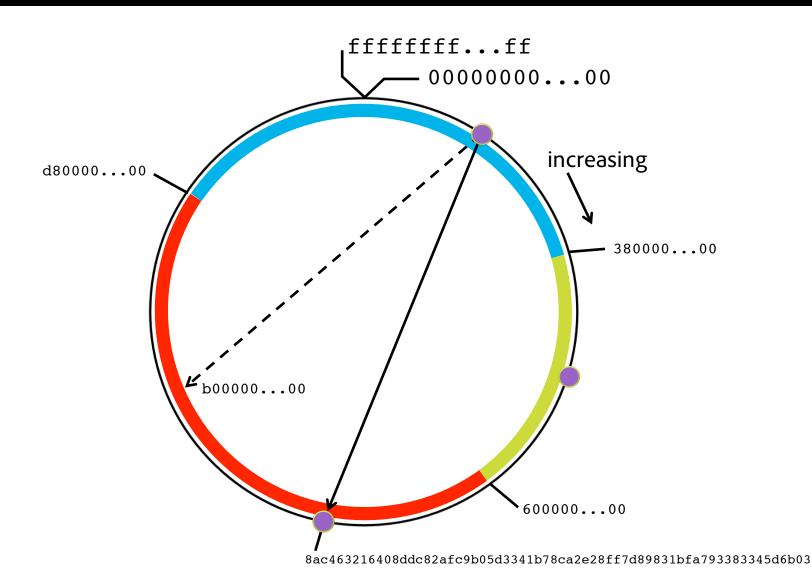
Send to Nearest



- NetGroup.sendToNearest() sends a message one hop to the neighbor (or self) numerically nearest to the target group address
 - EXACT peers are only nearest to their own group address
 - Each peer is nearest to its own group address
 - Recursive operation can route a message through the group hop-by-hop to the node that is actually closest
 - If target address is in self's local coverage, message is sent to self with fromLocal=true
 - Can build a DHT this way
 - Worst-case diameter is log₄ N
 - Often better

Send To Nearest





Send To Nearest



- Too many EXACT peers can also disrupt NEAREST routing
 - Ring-adjacent NEAREST peers must be neighbors with each other to calculate consistent non-overlapping local coverage ranges
 - A peer may over-estimate its local coverage, causing a message to not be routed to the actual nearest peer

Send To Neighbor



- NetGroup.sendToNeighbor() sends a message along the ring
 - NetGroupSendMode.NEXT_INCREASING, .NEXT_DECREASING
- Correct topology means self→NEXT_INCREASING→NEXT_DECREASING == self→NEXT_DECREASING→NEXT_INCREASING == self
 - Could traverse entire group hop-by-hop
- Incorrect topology means spurs in the ring
 - Potentially inconsistent recursive routing behavior
 - Usually caused by NAT or firewalls blocking a neighbor connection

Send To Neighbor



- Adjacency scales as O(1)
- Convenient when you need someone else to do something for you
 - Example: announce when you leave the group, since you can't do it
- When all members' receive modes are NEAREST,
 NetGroup.LocalCoverage.Notify event can indicate an adjacent neighbor has changed
 - localCoverageTo changing indicates NEXT_INCREASING is different
 - localCoverageFrom changing indicates NEXT_DECREASING is different

Send To Neighbor



DEMO

Send To All Neighbors



- NetGroup.sendToAllNeighbors()
 - Sends to all directly connected neighbors, not to every member of the group
 - That's what NetGroup.post() is for
 - Sends to all neighbors at the same time
 - Could jam network if you have a lot of neighbors

■ Demo shortly... ©

Small Groups



- Small groups are fully meshed
 - Each member is a neighbor with every other member
 - Diameter 1
 - 14 or fewer members guarantees* a full mesh (see topology slides)
 - NetGroup.sendToAllNeighbors() reaches all* members directly
 - Much lower latency than posting
 - Load not shared by other peers
 - Lower overhead than posting for non-trivial messages
 - Ideal for N:N or where latency is important (like Real Time games)

Small Groups



- Neighbor disconnect means the peer is really gone
- Smallest useful group is 2 members
 - As of Flash Player 11.0, can use 1:1 DIRECT_CONNECTIONS NetStream even in serverless mode
- NetGroup.sendToNearest() to exact address for easy 1:1 messaging

Large Groups



- Large groups are not fully meshed
 - Each member will not be a direct neighbor with every other member
 - Probability of stable full mesh drops rapidly above 14 members
 - Very unlikely above about N=20, probability 0 at N=33
- NetGroup.sendToAllNeighbors() doesn't reach all members
- Use posting for N:N communication or for infrequent short messages few:N
 - Higher latency, possibly seconds or more to reach all members of a large group
 - Members share load
 - Will reach all members even when NAT/firewalls break correct routing topology

Large and Small Groups



DEMO

Large Groups



- Useful as a lobby or discovery area, break-out to "small" groups
- Use NetGroup.sendToAllNeighbors() to reach O(log n) randomish subset
- Use multicast with NetStream.send() for continuous data from 1 (or a few) to all members
 - Possible jitter up to multicast window duration
 - Even a stable group takes tens of seconds to converge to low-latency "push" mode

Large Groups



- Recursive directed routing
 - Chances of NAT/firewall disruption of correct routing topology high
 - For DHT applications, use replication to increase chances of at least one working path
 - Routing diameter (worst case) log₄ N, usually better
- Gossip approach for efficient (but potentially slow) information diffusion
 - Distributed presence
 - Diffusion diameter O(log N)
- Flooding diameter may approach O(log log N)

All Groups



- Full transitive connectivity regardless of size with high probability
 - Posting, flooding, gossip, object replication, multicast highly likely to reach all members

What About Multicast?



- Multicast is good at 1 (or few) –to-all distribution
 - Builds efficient spanning trees through group for low-latency push mode
 - NetStream.send() works in a multicast stream
- Multicast can have high delay/jitter
 - Pull mode at startup while trees are building
 - Peers joining and leaving can disrupt trees, pull mode while they rebuild
 - Tree building requires data to be flowing

What About Multicast?



- NetStream.multicastWindowDuration
 - Limits latency and jitter
 - Lower delay may mean lower reliability (more missed messages)
 - Much lower window durations will probably also require lowering the multicastAvailabilityUpdatePeriod and multicastFetchPeriod
 - Will increase protocol overhead
 - Set parameters before NetStream.publish()
- N:N Streaming
 - Full mesh of DIRECT_CONNECTIONS NetStreams is more efficient



Q&A



