

P2P on the Flash Platform with RTMFP

Matthew Kaufman

Sr. Computer Scientist, Project Lead


Adobe



What is “Peer-to-Peer”?


- Point-to-point

- Live streaming
- Document delivery




- “Swarming”

- Large-file download (possibly progressive)




- Live Application-Level Multicast

- Broadcast (1 to many, some latency tolerable)
- Interactive (many to many, or 1 to many with feedback, low latency required)



- Distributed Data Storage

- DHT-like structures to form distributed database




- **Point-to-point**

- Live streaming
- Document delivery




- **“Swarming”**

- Large-file download (possibly progressive)




- **Live Application-Level Multicast**

- Broadcast (1 to many, some latency tolerable)
- Interactive (many to many, or 1 to many with feedback, low latency required)



- **Distributed Data Storage**

- DHT-like structures to form distributed database




- **Point-to-point**

- Live streaming
- Document delivery




- **“Swarming”**

- Large-file download (possibly progressive)




- **Live Application-Level Multicast**

- Broadcast (1 to many, some latency tolerable)
- Interactive (many to many, or 1 to many with feedback, low latency required)




- **Distributed Data Storage**

- DHT-like structures to form distributed database




- Join and participate in a self-organizing peer-to-peer overlay network
- Use that network for:
 - Directed Routing
 - Object Replication
 - Posting
 - Application-Level Multicast
- Use Native IP Multicast together with Application-Level Multicast
 - “Fusion”


Publishing to (a few) Subscribers through FMS




Publishing to (a few) Subscribers Peer-to-Peer




Publishing to (many) Subscribers through FMS



Publishing to (many) Subscribers using RTMFP Groups




Publishing to (many) Subscribers using RTMFP Groups




- 1 to 1 (1 to few)
 - Communicate via Server
 - RTMP (RTMPT, RTMPS)
 - RTMFP
 - Communicate Directly
 - RTMFP
- 1 to Many, Many to Many
 - Communicate via Server
 - Communicate using RTMFP Groups

Full Mesh of Peers






Full Mesh of Peers

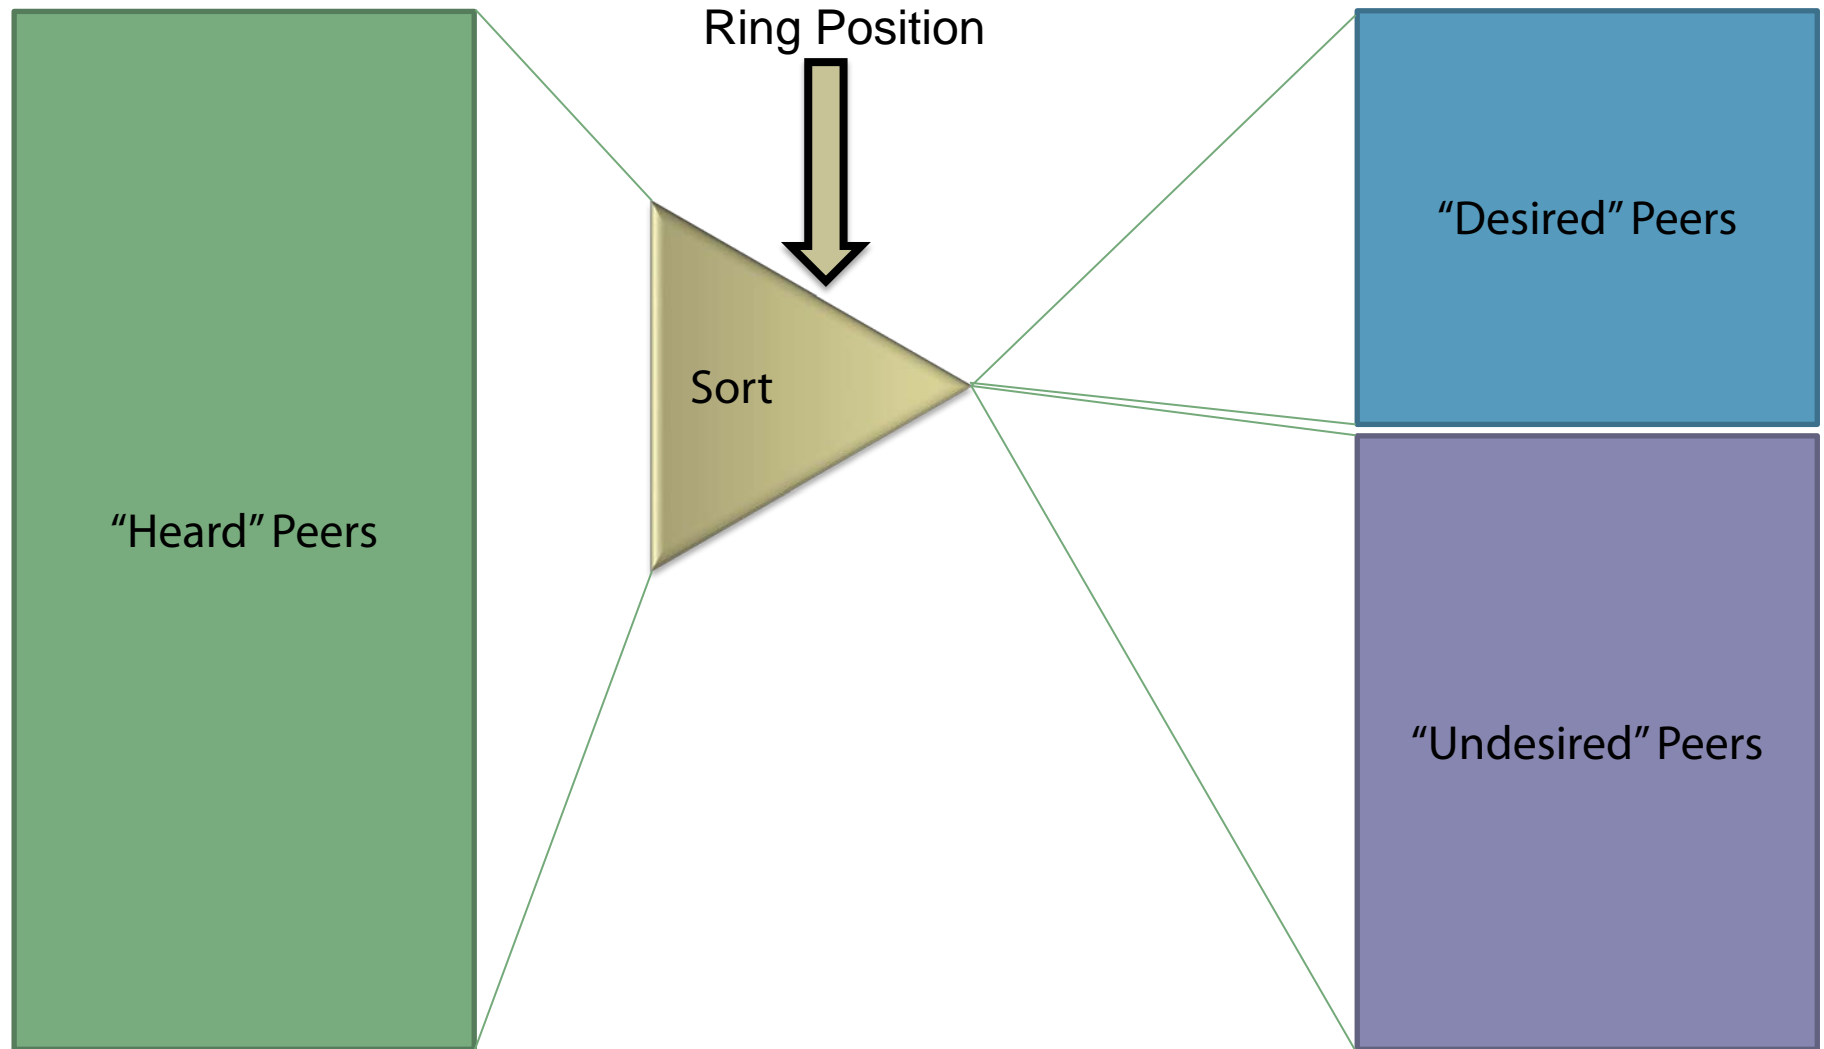


- Player needs to communicate with n peers efficiently
 - n might be very large (millions)
 - $O(n)$ at each peer ($O(n^2)$ overall) doesn't scale
- Solution: Overlay network
 - Topology optimized for multiple uses


- RTMFP is used as transport
 - NAT/Firewall traversal, Encryption, IP Address Mobility, Congestion Control, Partial Reliability, etc.
- Each peer has $O(\log n)$ connections
 - Typical: $2 \log_2 n + 13$
- Epidemic theory answers probability of connectedness
 - For $\log_2 n + c$ connections, probability is $e^{-e^{-c}}$
 - So c should be at least 10
- Topology management
 - decoupled from data transport
 - distributed to peers
- Controlled access
 - Ensures joiners have secret (the Groupspec) without exposing secret










- Initialize the “heard list” from one or more bootstrap nodes
- Repeat forever:
 - Sort the “heard list” to find what is best for self making a “best” list
 - Connect (or stay connected) to the nodes on the “best” list
 - Ask connected nodes that are no longer on the “best” list to disconnect
 - If a node asks you to disconnect, do so only if not in your “best” list
 - Every interval, pick one node at random that you are talking to
 - Sort the “heard list” to find what is best for them
 - Send them the result of that sort and ask them to do the same in return (“push-pull”)
 - Every time you hear about new nodes, merge them into the “heard list”
 - Delete nodes from “heard list” if their age limit is exceeded
- Resulting self-organizing topology depends entirely on sorting algorithm
- Continuous optimization through gossip, rate is higher than natural churn









- Nodes have a 256-bit numeric identifier ($0 \dots 2^{256} - 1$)
 - Number of dimensions = 1
- Sort chooses:
 - 6 closest (mod 2^{256}) in numeric space (3 closest on each side)
 - Numeric location + $1/2$, $1/4$, $1/8$, $1/16$, ...
 - 6 nodes with lowest latency
 - 1 random node (per round)
- Results in a fully-connected ring, plus:
 - Redundant neighbor connections
 - Shortcuts across the ring to reduce average hop count to less than $\log_2 n$
 - Additional nodes that are nearby (based on measured latency)
 - Random nodes sufficient to continue optimizing

- Groupspec
 - String starting with “G:” (e.g., “G:01010b.....”)
 - First part becomes the immutable “identity” of the Group
 - Certain options (permissions) must be in this part
 - If a peer tries to change, they are in a Group, but not this one
 - Note: Case-sensitive
- GroupSpecifier class makes them

■ Properties

- `routingEnabled`
- `multicastEnabled`
- `objectReplicationEnabled`
- `postingEnabled`
- `peerToPeerDisabled`
- `ipMulticastMemberUpdatesEnabled`
- `serverChannelEnabled`

■ Methods

- `GroupSpecifier(name:String)`
- `makeUnique()`
- `setPublishPassword()`
- `setPostingPassword()`
- `addBootstrapPeer()`
- `addIPMulticastAddress()`

■ Getting the Groupspec

- `groupspecWithAuthorizations()`
- `groupspecWithoutAuthorizations()`

- Ring Position
 - 256-bit numeric identifier
 - Formed from PeerID

- NetGroup
 - Management and statistics
 - All Group functionality except Multicast
- NetStream
 - Multicast
- More than one object can refer to the same Group
 - As long as in same NetConnection, overhead is shared

- **Properties**

- `estimatedMemberCount`
- `neighborCount`
- `info`
 - `NetGroupInfo` properties

- **Methods**

- `addNeighbor`
- `addMemberHint`
- `close`

- **NetStatusEvents**

- `NetGroup.Neighbor.Connect`
- `NetGroup.Neighbor.Disconnect`

- Connect

- Peer-Assisted Networking dialog will be displayed
 - When NetStream or NetGroup is constructed
 - Unless user remembers “allow” for your domain,
 - or `peerToPeerDisabled` is set in `GroupSpecifier`

- Use

- Directed Routing
- Object Replication
- Posting
- Multicast
 - Fusion

- Requires stable and correct topology to work well
- Works with ring positions
- API
 - `receiveMode`
 - `NetGroupReceiveMode.EXACT` (Initialized in this mode)
 - `NetGroupReceiveMode.NEAREST`
 - `sendToNearest`
 - `sendToNeighbor`
 - `NetGroupSendMode.NEXT_INCREASING`, `NetGroupSendMode.NEXT DECREASING`
 - `sendToAllNeighbors`
- At each hop, arrives in ActionScript as event
 - `NetGroup.SendTo.Notify`

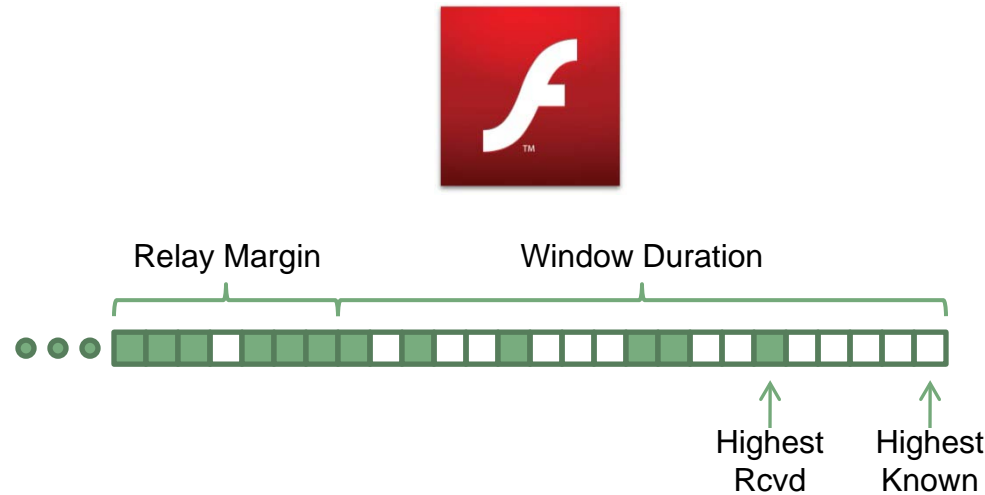
- Each hop can
 - Forward, Modify, Reply, or Discard
- Use cases:
 - Replying to specific node
 - Distributed Hash Table database
 - Research projects

- Objects are addressed by numeric Index
- ActionScript provides backing store
 - Easier to use files from AIR
- API
 - `addHaveObjects / removeHaveObjects`
 - `addWantObjects / removeWantObjects`
 - `writeRequestedObject / denyRequestedObject`
 - `NetGroup.Replication.Fetch.Result`
 - `NetGroup.Replication.Request`
- Use cases:
 - Workspace replication
 - Whiteboard
 - File transfer
 - ...


- One Object ID space per group
- Fully reliable delivery
 - All other RTMFP Groups data transports are best-effort
- Hints for moving around large amounts of data:
 - Break up into reasonable-sized objects (64k)
 - Consider using block 0 as descriptive object block
 - Consider making Group “name” be related to hash of descriptive object block

- Per-sender state (efficiency) tradeoff
 - Multicast is for a small # of senders @ high data rates
 - Posting is for a large # of senders @ low data rates
- API
 - `post()`
 - `NetGroup.Posting.Notify`
- Hints:
 - Messages must be unique to be considered new
 - Add a serial number
 - Keep messages small and relatively infrequent per sender
- Works without ActionScript involvement (unlike previous two modes)
- Use cases:
 - Text chat, sensor reporting, opinion polls


- Uses NetStream API
 - Construct with Groupspect instead of PeerID
- Designed for low latency
- Again, works without ActionScript involvement



- Multicast system attempts to get every block in the Window Duration
 - Once all blocks arrive for a message, message is processed
 - If any block of a message falls outside the Window, system no longer attempts to get any other block of that message
- Relay Margin provides extra time for in-flight requests to be fulfilled
- Both are in terms of time



- Send map of what blocks you have to one (or all) neighbor(s)
 - Neighbor waits a short time to see who can source each block





- Send map of what blocks you have to one (or all) neighbor(s)
 - Neighbor waits a short time to see who can source each block
- Neighbor picks you to send some of the blocks and sends request





- Send map of what blocks you have to one (or all) neighbor(s)
 - Neighbor waits a short time to see who can source each block
- Neighbor picks you to send some of the blocks and sends request
- Send requested blocks to neighbor
 - All transmissions are partially-reliable
 - Neighbor will ask someone else if request goes unfulfilled














































- Send map of what blocks you have to one (or all) neighbor(s)
 - Neighbor notices that data is flowing





- Neighbor picks candidate sources for each sequence number slice
- Neighbor sends mask describing which sequence numbers to push
 - Modulo size of mask (in above example every 3rd, 5th, and 6th)
 - Periodically tests other neighbors to check for lower latency
 - For each sequence number slice, keeps just quickest source





- Data is pushed immediately as it arrives for requested slices
- For each slice, there is a limit to the number of push clients served






















MAX
































- Uses NetStream API
 - Construct with Groupspec instead of PeerID
- Low latency
 - Push/pull
 - Pull
 - Ensure nodes get the data
 - Initial source
 - Back-fill of loss (Fusion or Push)
 - Provides data when spanning tree temporarily wrong
 - Push
 - Reduce latency
 - Multiple spanning trees
- >1 Sender/group
- >1 Stream/sender (streams have names)

■ Properties

- `multicastPushNeighborLimit`
- `multicastWindowDuration`
- `multicastRelayMarginDuration`
- `multicastAvailabilityUpdatePeriod`
- `multicastFetchPeriod`
- `multicastAvailabilitySendToAll`
- `multicastInfo`

■ Events

- `NetStream.MulticastStream.Reset`
- `NetGroup.MulticastStream.PublishNotify`
- `NetGroup.MulticastStream.UnpublishNotify`

- `sendDataBytesPerSecond`
- `sendControlBytesPerSecond`
- `receiveDataBytesPerSecond`
- `receiveControlBytesPerSecond`
- `bytesPushedToPeers`
- `fragmentsPushedToPeers`
- `bytesRequestedByPeers`
- `fragmentsRequestedByPeers`
- `bytesPushedFromPeers`
- `fragmentsPushedFromPeers`
- `bytesRequestedFromPeers`
- `fragmentsRequestedFromPeers`
- `sendControlBytesPerSecondToServer`
- `receiveDataBytesPerSecondFromServer`
- `bytesReceivedFromServer`
- `fragmentsReceivedFromServer`
- `receiveDataBytesPerSecondFromIPMulticast`
- `bytesReceivedFromIPMulticast`
- `fragmentsReceivedFromIPMulticast`

- Provide Native IP Multicast address(es) in Groupspec
- LAN Discovery
 - Increase odds of finding nearby (low-latency) peers
 - Serverless ad-hoc RTMFP Group using new `NetConnection.connect("rtmfp:");`

- Fuse native IP multicast with peer-to-peer Multicast
 - Peers can use pull to fill in missing data
 - Peers with native IP can feed to peers without
- Single Groupspec (stream description) can be played in either environment
- Turn off peer-to-peer and use for Native IP-only
- (Future) FMS is publisher
 - Flash Player cannot **publish** Native IP Multicast
- Security of Native IP Multicast traffic (Fusion and LAN Discovery)
 - Protected (and separated from other traffic) with Groupspec-derived HMAC
 - Encrypted with Groupspec-derived AES key

- Enabled in Groupspec
 - `GroupSpecifier.serverChannelEnabled`
- Status monitoring
- Multicast
 - Server seed
 - Server fill-in
- Bootstrapping
 - Server can tell Flash Player about other peers to add
- New Stratus feature!
 - Group auto-bootstrapping

- Partial Reliability Control at publishing peer
 - `NetStream.audioReliable`
 - `NetStream.videoReliable`
 - `NetStream.dataReliable`
- Peer-to-Peer Audio/Video Sample Access Control at publishing peer
 - `NetStream.audioSampleAccess`
 - `NetStream.videoSampleAccess`
- RTMFP NetConnection with no server
 - `NetConnection.connect("rtmfp:");`



CONNECT. DISCOVER. INSPIRE.